

# 10 Linear Regression Algorithms with Python

Linear Regression is the most straightforward algorithm in machine learning, it can be trained differently. In this notebook we will cover the following linear algorithms:

1. Linear Regression
2. Robust Regression
3. Ridge Regression
4. LASSO Regression
5. Elastic Net
6. Polynomial Regression
7. Stochastic Gradient Descent
8. Artificial Neural Networks
9. Random Forest Regressor
10. Support Vector Machine

## Data

We are going to use the `USA_Housing` dataset. Since house price is a continuous variable, this is a regression problem. The data contains the following columns:

- 'Avg. Area Income': Avg. The income of residents of the city house is located in.
- 'Avg. Area House Age': Avg Age of Houses in the same city
- 'Avg. Area Number of Rooms': Avg Number of Rooms for Houses in the same city
- 'Avg. Area Number of Bedrooms': Avg Number of Bedrooms for Houses in the same city
- 'Area Population': The population of city house is located in
- 'Price': Price that the house sold at
- 'Address': Address for the house

```
In [16]: !pip install -q hvplot
```

## Import Libraries

```
In [17]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import hvplot.pandas

%matplotlib inline

# sns.set_style("whitegrid")
# plt.style.use("fivethirtyeight")
```

## Check out the Data

```
In [18]: USAhousing = pd.read_csv('/input/usa-housing/USA_Housing.csv')
USAhousing.head()
```

```
Out[18]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386

```
In [19]: USAhousing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Avg. Area Income    5000 non-null   float64
 1   Avg. Area House Age 5000 non-null   float64
 2   Avg. Area Number of Rooms 5000 non-null   float64
 3   Avg. Area Number of Bedrooms 5000 non-null   float64
 4   Area Population     5000 non-null   float64
 5   Price               5000 non-null   float64
 6   Address             5000 non-null   object  
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

In [20]: USAhousing.describe()

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
<b>count</b>	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
<b>mean</b>	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
<b>std</b>	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
<b>min</b>	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
<b>25%</b>	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
<b>50%</b>	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
<b>75%</b>	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
<b>max</b>	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06

In [21]: USAhousing.columns

```
Out[21]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
       'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')
```

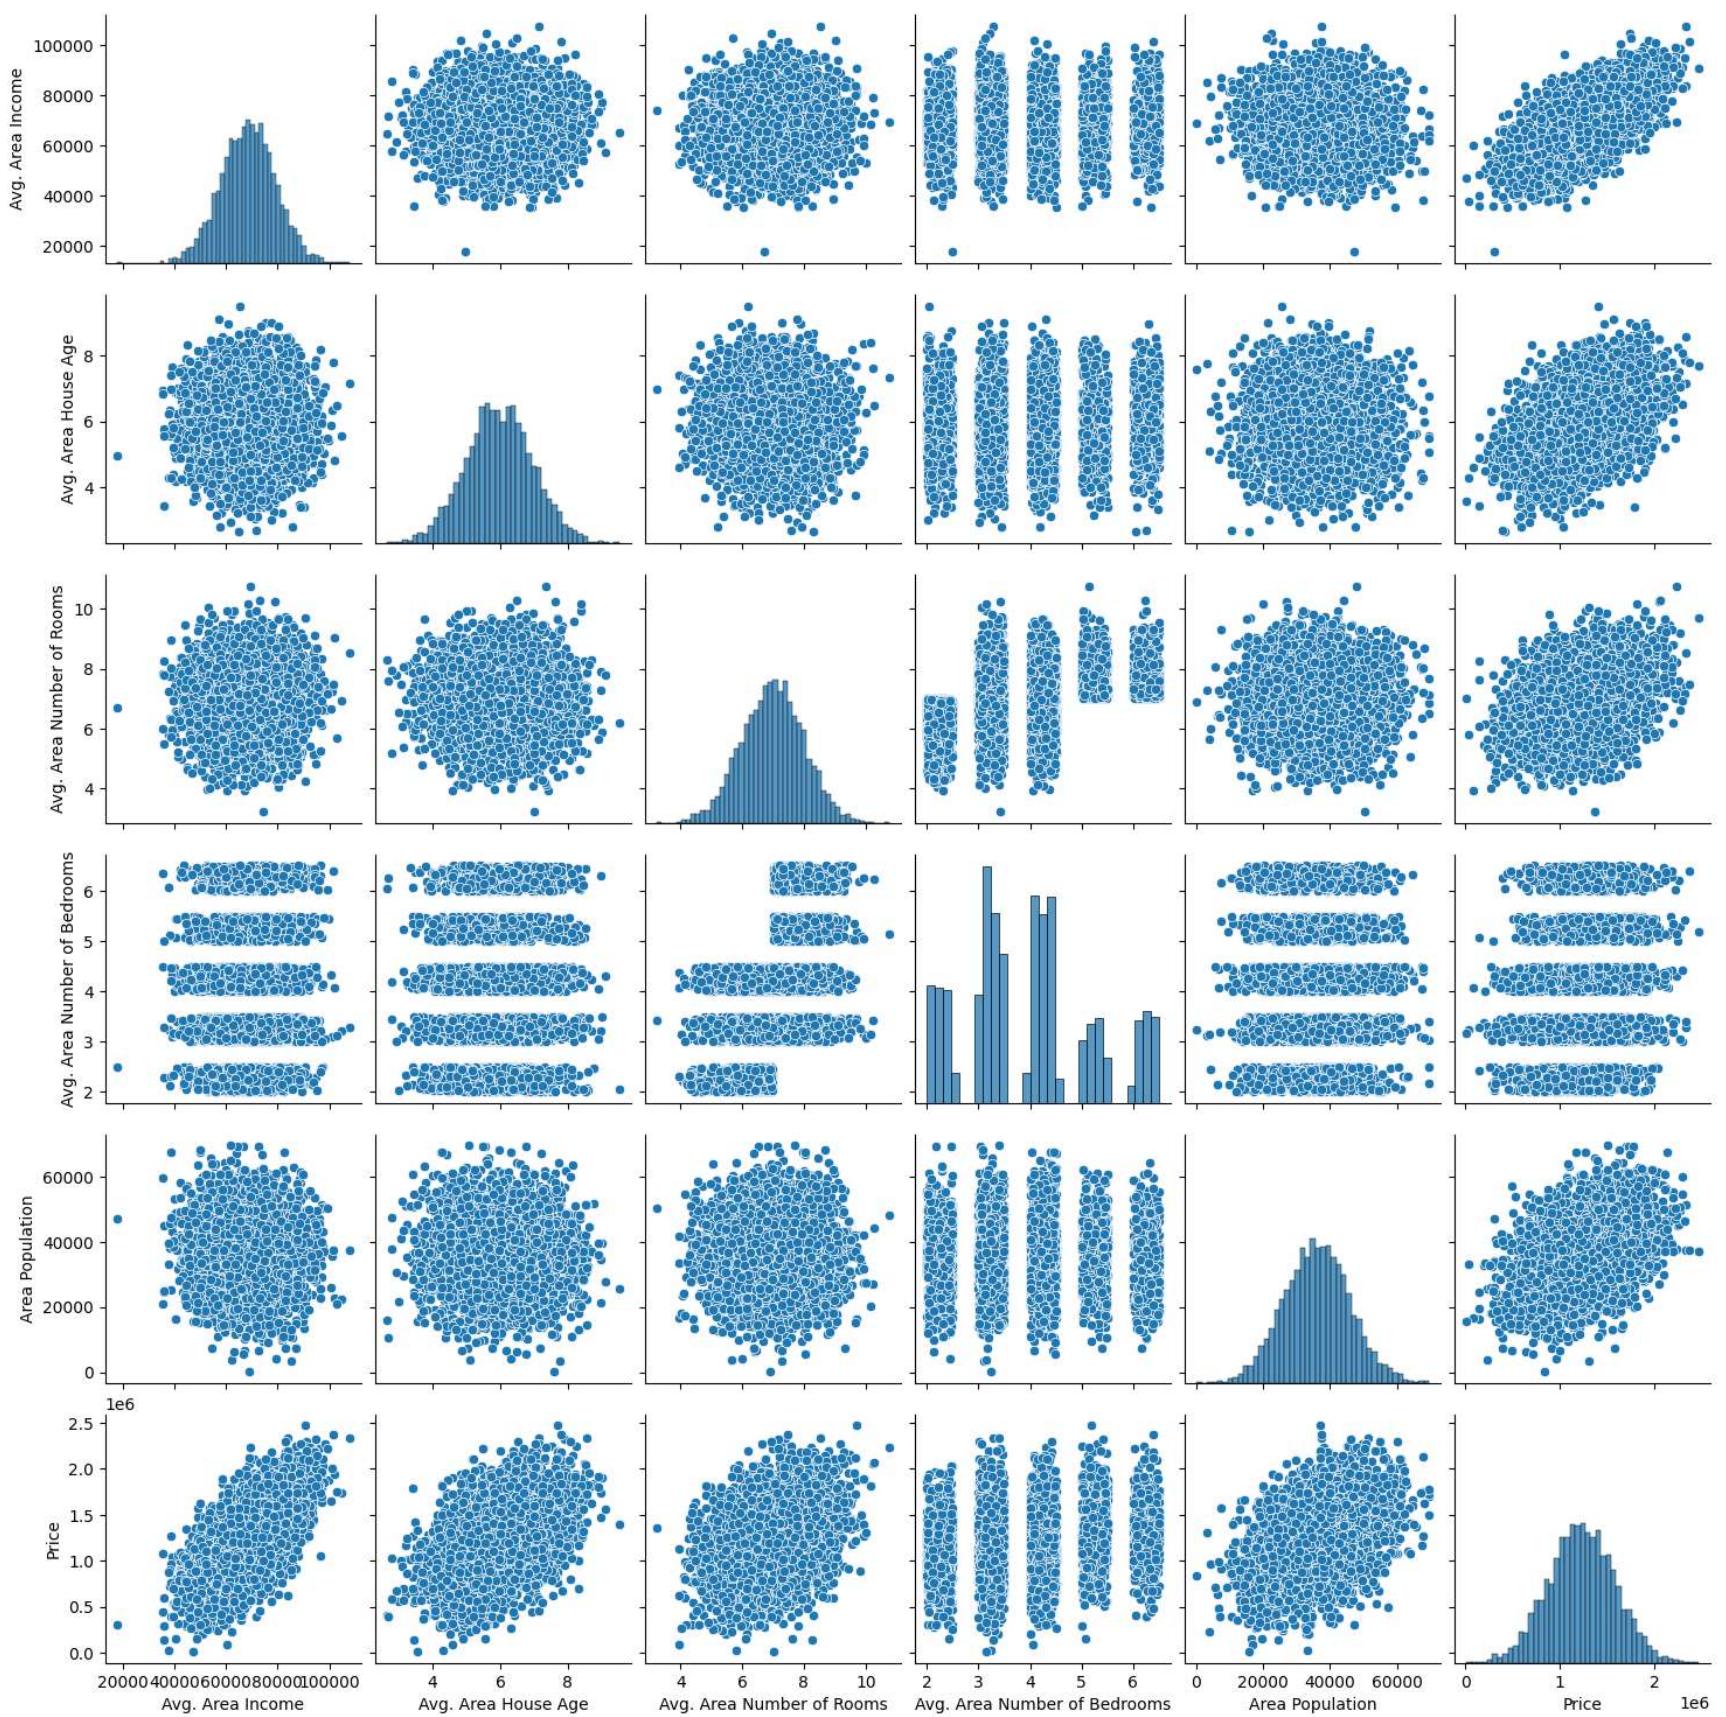
## Exploratory Data Analysis (EDA)

Let's create some simple plots to check out the data

In [22]: sns.pairplot(USAhousing)

```
/opt/conda/lib/python3.10/site-packages/seaborn/axisgrid.py:118: UserWarning: The figure layout has changed to tight
  self._figure.tight_layout(*args, **kwargs)
```

Out[22]: <seaborn.axisgrid.PairGrid at 0x7cdb0961aa10>



```
In [23]: USAhousing.hvplot.hist(by='Price', subplots=False, width=1000)
```

Out[23]:

```
In [24]: USAhousing.hvplot.hist("Price")
```

Out[24]:

```
In [25]: USAhousing.hvplot.scatter(x='Avg. Area House Age', y='Price')
```

Out[25]:

```
In [26]: USAhousing.hvplot.scatter(x='Avg. Area Income', y='Price')
```

Out[26]:

```
In [27]: USAhousing.columns
```

```
Out[27]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
       'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')
```

## Training a Linear Regression Model

We will need to first split up our data into an X array that contains the features to train on, and a y array with the target variable, in this case, the Price column. We will toss out the Address column because it only has text info that the linear regression model can't use.

### X and y arrays

```
In [28]: X = USAhousing[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
                     'Avg. Area Number of Bedrooms', 'Area Population']]
y = USAhousing['Price']
```

## Train Test Split

Now, let's split the data into a training set and a testing set. We will train our model on the training set and then use the test set to evaluate the model.

```
In [29]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [30]: from sklearn import metrics
from sklearn.model_selection import cross_val_score

def cross_val(model):
    pred = cross_val_score(model, X, y, cv=10)
    return pred.mean()

def print_evaluate(true, predicted):
    mae = metrics.mean_absolute_error(true, predicted)
    mse = metrics.mean_squared_error(true, predicted)
    rmse = np.sqrt(metrics.mean_squared_error(true, predicted))
    r2_square = metrics.r2_score(true, predicted)
    print('MAE:', mae)
    print('MSE:', mse)
    print('RMSE:', rmse)
    print('R2 Square', r2_square)
    print('_____')

def evaluate(true, predicted):
    mae = metrics.mean_absolute_error(true, predicted)
    mse = metrics.mean_squared_error(true, predicted)
    rmse = np.sqrt(metrics.mean_squared_error(true, predicted))
    r2_square = metrics.r2_score(true, predicted)
    return mae, mse, rmse, r2_square
```

## Preparing Data For Linear Regression

Try different preparations of your data using these heuristics and see what works best for your problem.

- **Linear Assumption.** Linear regression assumes a linear relationship between your input and output. It does not support any other type of relationship. While this may seem obvious, it's important to remember, especially when dealing with numerous attributes. In such cases, you might need to transform the data to establish a linear relationship, for instance, by applying a log transform for data exhibiting an exponential relationship.
- **Remove Noise.** Linear regression assumes that your input and output variables are not noisy. Consider using data cleaning operations that let you better expose and clarify the signal in your data. This is most important for the output variable and you want to remove outliers in the output variable ( $y$ ) if possible.
- **Remove Collinearity.** Linear regression will over-fit your data when you have highly correlated input variables. Consider calculating pairwise correlations for your input data and removing the most correlated.
- **Gaussian Distributions.** Linear regression will make more reliable predictions if your input and output variables have a Gaussian distribution. You may get some benefit using transforms (e.g. log or BoxCox) on your variables to make their distribution more Gaussian looking.
- **Rescale Inputs:** Linear regression will often make more reliable predictions if you rescale input variables using standardization or normalization.

```
In [31]: from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline

pipeline = Pipeline([
    ('std_scalar', StandardScaler())
])

X_train = pipeline.fit_transform(X_train)
X_test = pipeline.transform(X_test)
```

## Linear Regression

```
In [32]: from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(X_train,y_train)
```

```
Out[32]: ▾ LinearRegression
LinearRegression()
```

## Model Evaluation

Let's evaluate the model by checking out its coefficients and how we can interpret them.

```
In [33]: # print the intercept
print(lin_reg.intercept_)

1228219.1492415662
```

```
In [34]: coeff_df = pd.DataFrame(lin_reg.coef_, X.columns, columns=['Coefficient'])

coeff_df
```

```
Out[34]:
```

	Coefficient
<b>Avg. Area Income</b>	232679.724643
<b>Avg. Area House Age</b>	163841.046593
<b>Avg. Area Number of Rooms</b>	121110.555478
<b>Avg. Area Number of Bedrooms</b>	2892.815119
<b>Area Population</b>	151252.342377

Interpreting the coefficients:

- Holding all other features fixed, a 1 unit increase in **Avg. Area Income** is associated with an **increase of \$23.26**.
- Holding all other features fixed, a 1 unit increase in **Avg. Area House Age** is associated with an **increase of \$163841**.
- Holding all other features fixed, a 1 unit increase in **Avg. Area Number of Rooms** is associated with an **increase of \$121110**.
- Holding all other features fixed, a 1 unit increase in **Avg. Area Number of Bedrooms** is associated with an **increase of \$2892.81**.
- Holding all other features fixed, a 1 unit increase in **Area Population** is associated with an **increase of \$15.12**.

Does this make sense? Probably not because I made up this data.

## Predictions from our Model

Let's grab predictions off our test set and see how well it did

```
In [35]: pred = lin_reg.predict(X_test)
```

```
In [36]: pd.DataFrame({'True Values': y_test, 'Predicted Values': pred}).hvplot.scatter(x='True Values', y='Predicted Values')
```

```
Out[36]:
```

## Residual Histogram

```
In [37]: pd.DataFrame({'Error Values': (y_test - pred)}).hvplot.kde()
```

```
Out[37]:
```

## Regression Evaluation Metrics

Here are three common evaluation metrics for regression problems:

- **Mean Absolute Error** (MAE) is the mean of the absolute value of the errors:

$$\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

- **Mean Squared Error** (MSE) is the mean of the squared errors:

$$\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- **Root Mean Squared Error** (RMSE) is the square root of the mean of the squared errors:

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

Comparing these metrics:

- **MAE** is the easiest to understand, because it's the average error.
- **MSE** is more popular than MAE, because MSE "punishes" larger errors, which tends to be useful in the real world.
- **RMSE** is even more popular than MSE, because RMSE is interpretable in the "y" units.

All of these are **loss functions**, because we want to minimize them.

```
In [38]: test_pred = lin_reg.predict(X_test)
train_pred = lin_reg.predict(X_train)
```

```

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df = pd.DataFrame(data=[["Linear Regression", *evaluate(y_test, test_pred), cross_val(LinearRegression())]], columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])

```

Test set evaluation:

---

MAE: 81135.56609336878  
MSE: 10068422551.40088  
RMSE: 100341.52954485436  
R2 Square 0.9146818498754016

Train set evaluation:

---

MAE: 81480.49973174892  
MSE: 10287043161.197224  
RMSE: 101425.06180031257  
R2 Square 0.9192986579075526

---

## Robust Regression

Robust regression is a form of regression analysis designed to overcome some limitations of traditional parametric and non-parametric methods. Robust regression methods are designed to be not overly affected by violations of assumptions by the underlying data-generating process.

One instance in which robust estimation should be considered is when there is a strong suspicion of **heteroscedasticity**.

A common situation in which robust estimation is used occurs when the data contain outliers. In the presence of outliers that do not come from the same data-generating process as the rest of the data, least squares estimation is inefficient and can be biased. Because the least squares predictions are dragged towards the outliers, and because the variance of the estimates is artificially inflated, the result is that outliers can be masked. (In many situations, including some areas of geostatistics and medical statistics, it is precisely the outliers that are of interest.)

## Random Sample Consensus - RANSAC

Random sample consensus (**RANSAC**) is an iterative method to estimate parameters of a mathematical model from a set of observed data that contains outliers, when outliers are to be accorded no influence on the values of the estimates. Therefore, it also can be interpreted as an outlier detection method.

A basic assumption is that the data consists of "inliers", i.e., data whose distribution can be explained by some set of model parameters, though may be subject to noise, and "outliers" which are data that do not fit the model. The outliers can come, for example, from extreme values of the noise or from erroneous measurements or incorrect hypotheses about the interpretation of data. RANSAC also assumes that, given a (usually small) set of inliers, there exists a procedure which can estimate the parameters of a model that optimally explains or fits this data.

In [45]:

```

from sklearn.linear_model import RANSACRegressor

model = RANSACRegressor(base_estimator=LinearRegression(), max_trials=100)
model.fit(X_train, y_train)

test_pred = model.predict(X_test)
train_pred = model.predict(X_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Robust Regression", *evaluate(y_test, test_pred), cross_val(RANSACRegressor())]], columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])

```

/opt/conda/lib/python3.10/site-packages/sklearn/linear\_model/\_ransac.py:343: FutureWarning: `base\_estimator` was renamed to `estimator` in version 1.1 and will be removed in 1.3.  
warnings.warn(  
Test set evaluation:

---

MAE: 82029.41475274497  
MSE: 10283624826.83934  
RMSE: 101408.20887304607  
R2 Square 0.9128582613292039

---

=====  
Train set evaluation:

---

MAE: 82712.18012372882  
MSE: 10599060832.325783  
RMSE: 102951.74030741677  
R2 Square 0.9168508947921401

## Ridge Regression

Ridge regression addresses some of the problems of **Ordinary Least Squares** by imposing a penalty on the size of coefficients. The ridge coefficients minimize a penalized residual sum of squares,

$$\min_w \|Xw - y\|_2^2 + \alpha \|w\|_2^2$$

$\alpha \geq 0$  is a complexity parameter that controls the amount of shrinkage: the larger the value of  $\alpha$ , the greater the amount of shrinkage and thus the coefficients become more robust to collinearity.

Ridge regression is an L2 penalized model. Add the squared sum of the weights to the least-squares cost function.

```
In [46]: from sklearn.linear_model import Ridge

model = Ridge(alpha=100, solver='cholesky', tol=0.0001, random_state=42)
model.fit(X_train, y_train)
pred = model.predict(X_test)

test_pred = model.predict(X_test)
train_pred = model.predict(X_train)

print('Test set evaluation:\n' + '-----')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n' + '-----')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Ridge Regression", *evaluate(y_test, test_pred), cross_val(Ridge())]], columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])
```

Test set evaluation:

MAE: 81428.64835535336  
MSE: 10153269900.892612  
RMSE: 100763.43533689497  
R2 Square 0.9139628674464607

=====

Train set evaluation:

MAE: 81972.39058585512  
MSE: 10382929615.14346  
RMSE: 101896.66145239233  
R2 Square 0.9185464334441484

## LASSO Regression

A linear model that estimates sparse coefficients.

Mathematically, it consists of a linear model trained with  $\ell_1$  prior as regularizer. The objective function to minimize is:

$$\min_w \frac{1}{2n_{samples}} \|Xw - y\|_2^2 + \alpha \|w\|_1$$

The lasso estimate thus solves the minimization of the least-squares penalty with  $\alpha \|w\|_1$  added, where  $\alpha$  is a constant and  $\|w\|_1$  is the  $\ell_1$ -norm of the parameter vector.

```
In [48]: from sklearn.linear_model import Lasso

model = Lasso(alpha=0.1,
              precompute=True,
#               warm_start=True,
              positive=True,
              selection='random',
              random_state=42)
model.fit(X_train, y_train)

test_pred = model.predict(X_test)
train_pred = model.predict(X_train)

print('Test set evaluation:\n' + '-----')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n' + '-----')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Lasso Regression", *evaluate(y_test, test_pred), cross_val(Lasso())]], columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])
```

Test set evaluation:

---

MAE: 81135.69851726218  
MSE: 10068453390.364521  
RMSE: 100341.68321472648  
R2 Square 0.914681588551116

---

Train set evaluation:

---

MAE: 81480.63002185506  
MSE: 10287043196.634295  
RMSE: 101425.0619750084  
R2 Square 0.9192986576295505

---

## Elastic Net

A linear regression model trained with L1 and L2 prior as regularizer. This combination allows for learning a sparse model where few of the weights are non-zero like Lasso, while still maintaining the regularization properties of Ridge. Elastic-net is useful when there are multiple features which are correlated with one another. Lasso is likely to pick one of these at random, while elastic-net is likely to pick both. A practical advantage of trading-off between Lasso and Ridge is it allows Elastic-Net to inherit some of Ridge's stability under rotation. The objective function to minimize is in this case

$$\min_w \frac{1}{2n_{samples}} \|Xw - y\|_2^2 + \alpha\rho\|w\|_1 + \frac{\alpha(1-\rho)}{2} \|w\|_2^2$$

```
In [49]: from sklearn.linear_model import ElasticNet

model = ElasticNet(alpha=0.1, l1_ratio=0.9, selection='random', random_state=42)
model.fit(X_train, y_train)

test_pred = model.predict(X_test)
train_pred = model.predict(X_train)

print('Test set evaluation:\n' + '-----')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n' + '-----')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Elastic Net Regression", *evaluate(y_test, test_pred), cross_val(ElasticNet())]], columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', "Cross Validation"])
```

Test set evaluation:

---

MAE: 81184.43147330944  
MSE: 10078050168.470106  
RMSE: 100389.49232100991  
R2 Square 0.9146002670381437

---

Train set evaluation:

---

MAE: 81577.88831531754  
MSE: 10299274948.101461  
RMSE: 101485.34351373829  
R2 Square 0.9192027001474953

---

## Polynomial Regression

One common pattern within machine learning is to use linear models trained on nonlinear functions of the data. This approach maintains the generally fast performance of linear methods, while allowing them to fit a much wider range of data.

For example, a simple linear regression can be extended by constructing polynomial features from the coefficients. In the standard linear regression case, you might have a model that looks like this for two-dimensional data:

$$\hat{y}(w, x) = w_0 + w_1x_1 + w_2x_2$$

If we want to fit a paraboloid to the data instead of a plane, we can combine the features in second-order polynomials, so that the model looks like this:

$$\hat{y}(w, x) = w_0 + w_1x_1 + w_2x_2 + w_3x_1x_2 + w_4x_1^2 + w_5x_2^2$$

The (sometimes surprising) observation is that this is still a linear model: to see this, imagine creating a new variable

$$z = [x_1, x_2, x_1x_2, x_1^2, x_2^2]$$

With this re-labeling of the data, our problem can be written

$$\hat{y}(w, x) = w_0 + w_1z_1 + w_2z_2 + w_3z_3 + w_4z_4 + w_5z_5$$

We see that the resulting polynomial regression is in the same class of linear models we'd considered above (i.e. the model is linear in w) and can be solved by the same techniques. By considering linear fits within a higher-dimensional

space built with these basis functions, the model has the flexibility to fit a much broader range of data.

```
In [52]: from sklearn.preprocessing import PolynomialFeatures

poly_reg = PolynomialFeatures(degree=2)

X_train_2_d = poly_reg.fit_transform(X_train)
X_test_2_d = poly_reg.transform(X_test)

lin_reg = LinearRegression()
lin_reg.fit(X_train_2_d,y_train)

test_pred = lin_reg.predict(X_test_2_d)
train_pred = lin_reg.predict(X_train_2_d)

print('Test set evaluation:\n')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Polynomial Regression", *evaluate(y_test, test_pred), 0]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])
```

Test set evaluation:

---

```
MAE: 81174.51844119694
MSE: 10081983997.620693
RMSE: 100409.08324260656
R2 Square 0.914566932419506
```

---

Train set evaluation:

---

```
MAE: 81363.0618562117
MSE: 10266487151.007816
RMSE: 101323.67517519198
R2 Square 0.9194599187853729
```

---

## Stochastic Gradient Descent

Gradient Descent is a very generic optimization algorithm capable of finding optimal solutions to a wide range of problems. The general idea of Gradient Descent is to tweak parameters iteratively in order to minimize a cost function. Gradient Descent measures the local gradient of the error function with regards to the parameters vector, and it goes in the direction of descending gradient. Once the gradient is zero, you have reached a minimum.

```
In [53]: from sklearn.linear_model import SGDRegressor

sgd_reg = SGDRegressor(n_iter_no_change=250, penalty=None, eta0=0.0001, max_iter=1000000)
sgd_reg.fit(X_train, y_train)

test_pred = sgd_reg.predict(X_test)
train_pred = sgd_reg.predict(X_train)

print('Test set evaluation:\n')
print_evaluate(y_test, test_pred)
print('=====')
print('Train set evaluation:\n')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Stochastic Gradient Descent", *evaluate(y_test, test_pred), 0]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])
```

Test set evaluation:

---

```
MAE: 81135.56653706348
MSE: 10068422540.702381
RMSE: 100341.52949154393
R2 Square 0.914681849966059
```

---

Train set evaluation:

---

```
MAE: 81480.49988440325
MSE: 10287043161.199207
RMSE: 101425.06180032235
R2 Square 0.919298657907537
```

---

## Artificial Neural Network

```
In [54]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Input, Dense, Activation, Dropout
from tensorflow.keras.optimizers import Adam

X_train = np.array(X_train)
X_test = np.array(X_test)
y_train = np.array(y_train)
y_test = np.array(y_test)
```

```
model = Sequential()

model.add(Dense(X_train.shape[1], activation='relu'))
model.add(Dense(32, activation='relu'))
# model.add(Dropout(0.2))

model.add(Dense(64, activation='relu'))
# model.add(Dropout(0.2))

model.add(Dense(128, activation='relu'))
# model.add(Dropout(0.2))

model.add(Dense(512, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(1))

model.compile(optimizer=Adam(0.00001), loss='mse')

r = model.fit(X_train, y_train,
               validation_data=(X_test,y_test),
               batch_size=1,
               epochs=100)
```

Epoch 1/100  
3500/3500 [=====] - 17s 3ms/step - loss: 1635959308288.0000 - val\_loss: 1658109558784.0000  
Epoch 2/100  
3500/3500 [=====] - 11s 3ms/step - loss: 1635207741440.0000 - val\_loss: 1656222253056.0000  
Epoch 3/100  
3500/3500 [=====] - 11s 3ms/step - loss: 1630735171584.0000 - val\_loss: 1647849373696.0000  
Epoch 4/100  
3500/3500 [=====] - 11s 3ms/step - loss: 1615667134464.0000 - val\_loss: 1623587160064.0000  
Epoch 5/100  
3500/3500 [=====] - 11s 3ms/step - loss: 1578034659328.0000 - val\_loss: 1568294567936.0000  
Epoch 6/100  
3500/3500 [=====] - 11s 3ms/step - loss: 1500520251392.0000 - val\_loss: 1461720580096.0000  
Epoch 7/100  
3500/3500 [=====] - 11s 3ms/step - loss: 1362529615872.0000 - val\_loss: 1283692560384.0000  
Epoch 8/100  
3500/3500 [=====] - 12s 3ms/step - loss: 1151140757504.0000 - val\_loss: 1031173898240.0000  
Epoch 9/100  
3500/3500 [=====] - 11s 3ms/step - loss: 884213612544.0000 - val\_loss: 744122023936.0000  
Epoch 10/100  
3500/3500 [=====] - 11s 3ms/step - loss: 632530993152.0000 - val\_loss: 522820157440.0000  
Epoch 11/100  
3500/3500 [=====] - 12s 3ms/step - loss: 476224061440.0000 - val\_loss: 410893287424.0000  
Epoch 12/100  
3500/3500 [=====] - 11s 3ms/step - loss: 408242946048.0000 - val\_loss: 364421283840.0000  
Epoch 13/100  
3500/3500 [=====] - 12s 3ms/step - loss: 375205691392.0000 - val\_loss: 337152376832.0000  
Epoch 14/100  
3500/3500 [=====] - 11s 3ms/step - loss: 350786945024.0000 - val\_loss: 315824963584.0000  
Epoch 15/100  
3500/3500 [=====] - 11s 3ms/step - loss: 329439346688.0000 - val\_loss: 296637693952.0000  
Epoch 16/100  
3500/3500 [=====] - 11s 3ms/step - loss: 310290513920.0000 - val\_loss: 278684532736.0000  
Epoch 17/100  
3500/3500 [=====] - 11s 3ms/step - loss: 291863527424.0000 - val\_loss: 261565693952.0000  
Epoch 18/100  
3500/3500 [=====] - 11s 3ms/step - loss: 272524886016.0000 - val\_loss: 245535424512.0000  
Epoch 19/100  
3500/3500 [=====] - 12s 3ms/step - loss: 255604457472.0000 - val\_loss: 230537510912.0000  
Epoch 20/100  
3500/3500 [=====] - 11s 3ms/step - loss: 239092252672.0000 - val\_loss: 216159371264.0000  
Epoch 21/100  
3500/3500 [=====] - 11s 3ms/step - loss: 223985336320.0000 - val\_loss: 202297655296.0000  
Epoch 22/100  
3500/3500 [=====] - 12s 3ms/step - loss: 209053646848.0000 - val\_loss: 189402415104.0000  
Epoch 23/100  
3500/3500 [=====] - 11s 3ms/step - loss: 195379871744.0000 - val\_loss: 177145479168.0000  
Epoch 24/100  
3500/3500 [=====] - 12s 3ms/step - loss: 180810350592.0000 - val\_loss: 165425446912.0000  
Epoch 25/100  
3500/3500 [=====] - 11s 3ms/step - loss: 169834512384.0000 - val\_loss: 154421116928.0000  
Epoch 26/100  
3500/3500 [=====] - 11s 3ms/step - loss: 158785880064.0000 - val\_loss: 144324870144.0000  
Epoch 27/100  
3500/3500 [=====] - 12s 3ms/step - loss: 146980306944.0000 - val\_loss: 134498590720.0000  
Epoch 28/100  
3500/3500 [=====] - 11s 3ms/step - loss: 136268660736.0000 - val\_loss: 125185179648.0000  
Epoch 29/100  
3500/3500 [=====] - 11s 3ms/step - loss: 126436106240.0000 - val\_loss: 116454137856.0000  
Epoch 30/100  
3500/3500 [=====] - 12s 3ms/step - loss: 117105737728.0000 - val\_loss: 108315385856.0000  
Epoch 31/100  
3500/3500 [=====] - 11s 3ms/step - loss: 109463683072.0000 - val\_loss: 100691501056.0000  
Epoch 32/100  
3500/3500 [=====] - 11s 3ms/step - loss: 101181554688.0000 - val\_loss: 93567795200.0000  
Epoch 33/100  
3500/3500 [=====] - 12s 3ms/step - loss: 93273227264.0000 - val\_loss: 86965608448.0000  
Epoch 34/100  
3500/3500 [=====] - 11s 3ms/step - loss: 86881157120.0000 - val\_loss: 80709943296.0000  
Epoch 35/100  
3500/3500 [=====] - 12s 3ms/step - loss: 80447799296.0000 - val\_loss: 74884964352.0000  
Epoch 36/100  
3500/3500 [=====] - 11s 3ms/step - loss: 74682400768.0000 - val\_loss: 69521555456.0000  
Epoch 37/100  
3500/3500 [=====] - 11s 3ms/step - loss: 69652742144.0000 - val\_loss: 64572534784.0000  
Epoch 38/100  
3500/3500 [=====] - 12s 3ms/step - loss: 64600694784.0000 - val\_loss: 60053512192.0000  
Epoch 39/100  
3500/3500 [=====] - 11s 3ms/step - loss: 59956658176.0000 - val\_loss: 55881830400.0000  
Epoch 40/100  
3500/3500 [=====] - 11s 3ms/step - loss: 55683854336.0000 - val\_loss: 52117303296.0000  
Epoch 41/100  
3500/3500 [=====] - 12s 3ms/step - loss: 52298625024.0000 - val\_loss: 48704151552.0000  
Epoch 42/100  
3500/3500 [=====] - 11s 3ms/step - loss: 48704208896.0000 - val\_loss: 45662937088.0000  
Epoch 43/100  
3500/3500 [=====] - 12s 3ms/step - loss: 45889564672.0000 - val\_loss: 42953994240.0000  
Epoch 44/100  
3500/3500 [=====] - 11s 3ms/step - loss: 43336740864.0000 - val\_loss: 40526344192.0000  
Epoch 45/100  
3500/3500 [=====] - 11s 3ms/step - loss: 41343074304.0000 - val\_loss: 38329667584.0000  
Epoch 46/100  
3500/3500 [=====] - 12s 3ms/step - loss: 39019712512.0000 - val\_loss: 36401291264.0000  
Epoch 47/100  
3500/3500 [=====] - 11s 3ms/step - loss: 37273694208.0000 - val\_loss: 34713907200.0000  
Epoch 48/100

3500/3500 [=====] - 11s 3ms/step - loss: 35525615616.0000 - val\_loss: 33283368960.0000  
Epoch 49/100  
3500/3500 [=====] - 12s 3ms/step - loss: 34034669568.0000 - val\_loss: 31990720512.0000  
Epoch 50/100  
3500/3500 [=====] - 11s 3ms/step - loss: 32745592832.0000 - val\_loss: 30802511872.0000  
Epoch 51/100  
3500/3500 [=====] - 12s 3ms/step - loss: 31757709312.0000 - val\_loss: 29763026944.0000  
Epoch 52/100  
3500/3500 [=====] - 12s 3ms/step - loss: 30725447680.0000 - val\_loss: 28870731776.0000  
Epoch 53/100  
3500/3500 [=====] - 11s 3ms/step - loss: 30301433856.0000 - val\_loss: 28002928640.0000  
Epoch 54/100  
3500/3500 [=====] - 12s 3ms/step - loss: 29413881856.0000 - val\_loss: 27256612864.0000  
Epoch 55/100  
3500/3500 [=====] - 11s 3ms/step - loss: 28519575552.0000 - val\_loss: 26560821248.0000  
Epoch 56/100  
3500/3500 [=====] - 11s 3ms/step - loss: 27828357120.0000 - val\_loss: 25949415424.0000  
Epoch 57/100  
3500/3500 [=====] - 12s 3ms/step - loss: 27362172928.0000 - val\_loss: 25376983040.0000  
Epoch 58/100  
3500/3500 [=====] - 11s 3ms/step - loss: 26997155840.0000 - val\_loss: 24801363968.0000  
Epoch 59/100  
3500/3500 [=====] - 11s 3ms/step - loss: 26366113792.0000 - val\_loss: 24315068416.0000  
Epoch 60/100  
3500/3500 [=====] - 12s 3ms/step - loss: 25624805376.0000 - val\_loss: 23842801664.0000  
Epoch 61/100  
3500/3500 [=====] - 12s 3ms/step - loss: 25158729728.0000 - val\_loss: 23431553024.0000  
Epoch 62/100  
3500/3500 [=====] - 12s 3ms/step - loss: 24828473344.0000 - val\_loss: 22952957952.0000  
Epoch 63/100  
3500/3500 [=====] - 11s 3ms/step - loss: 24133113856.0000 - val\_loss: 22608713728.0000  
Epoch 64/100  
3500/3500 [=====] - 11s 3ms/step - loss: 24152397824.0000 - val\_loss: 22187180032.0000  
Epoch 65/100  
3500/3500 [=====] - 12s 3ms/step - loss: 23455307776.0000 - val\_loss: 21802670080.0000  
Epoch 66/100  
3500/3500 [=====] - 11s 3ms/step - loss: 22972102656.0000 - val\_loss: 21466863616.0000  
Epoch 67/100  
3500/3500 [=====] - 12s 3ms/step - loss: 22910437376.0000 - val\_loss: 21139642368.0000  
Epoch 68/100  
3500/3500 [=====] - 12s 3ms/step - loss: 22670950400.0000 - val\_loss: 20811978752.0000  
Epoch 69/100  
3500/3500 [=====] - 12s 3ms/step - loss: 22494793728.0000 - val\_loss: 20500830208.0000  
Epoch 70/100  
3500/3500 [=====] - 12s 3ms/step - loss: 22153525248.0000 - val\_loss: 20192782336.0000  
Epoch 71/100  
3500/3500 [=====] - 12s 3ms/step - loss: 22092113920.0000 - val\_loss: 19905064960.0000  
Epoch 72/100  
3500/3500 [=====] - 12s 3ms/step - loss: 21601464320.0000 - val\_loss: 19627452416.0000  
Epoch 73/100  
3500/3500 [=====] - 12s 3ms/step - loss: 21141583872.0000 - val\_loss: 19381852160.0000  
Epoch 74/100  
3500/3500 [=====] - 11s 3ms/step - loss: 20818229248.0000 - val\_loss: 19158253568.0000  
Epoch 75/100  
3500/3500 [=====] - 12s 3ms/step - loss: 20810127360.0000 - val\_loss: 18913705984.0000  
Epoch 76/100  
3500/3500 [=====] - 12s 3ms/step - loss: 20644577280.0000 - val\_loss: 18656931840.0000  
Epoch 77/100  
3500/3500 [=====] - 12s 3ms/step - loss: 20164083712.0000 - val\_loss: 18444519424.0000  
Epoch 78/100  
3500/3500 [=====] - 12s 3ms/step - loss: 19818045440.0000 - val\_loss: 18173810688.0000  
Epoch 79/100  
3500/3500 [=====] - 12s 3ms/step - loss: 19892738048.0000 - val\_loss: 17975521280.0000  
Epoch 80/100  
3500/3500 [=====] - 11s 3ms/step - loss: 19654350848.0000 - val\_loss: 17776916480.0000  
Epoch 81/100  
3500/3500 [=====] - 12s 3ms/step - loss: 19405084672.0000 - val\_loss: 17564684288.0000  
Epoch 82/100  
3500/3500 [=====] - 12s 3ms/step - loss: 19170838528.0000 - val\_loss: 17357641728.0000  
Epoch 83/100  
3500/3500 [=====] - 12s 3ms/step - loss: 19056832512.0000 - val\_loss: 17147287552.0000  
Epoch 84/100  
3500/3500 [=====] - 12s 3ms/step - loss: 18851735552.0000 - val\_loss: 16961145856.0000  
Epoch 85/100  
3500/3500 [=====] - 12s 3ms/step - loss: 18791458816.0000 - val\_loss: 16794759168.0000  
Epoch 86/100  
3500/3500 [=====] - 12s 3ms/step - loss: 18444681216.0000 - val\_loss: 16653736960.0000  
Epoch 87/100  
3500/3500 [=====] - 12s 3ms/step - loss: 18259365888.0000 - val\_loss: 16436895744.0000  
Epoch 88/100  
3500/3500 [=====] - 12s 3ms/step - loss: 18135013376.0000 - val\_loss: 16301491200.0000  
Epoch 89/100  
3500/3500 [=====] - 12s 3ms/step - loss: 18294607872.0000 - val\_loss: 16206785536.0000  
Epoch 90/100  
3500/3500 [=====] - 12s 3ms/step - loss: 17888196608.0000 - val\_loss: 15975628800.0000  
Epoch 91/100  
3500/3500 [=====] - 12s 3ms/step - loss: 17678673920.0000 - val\_loss: 15897917440.0000  
Epoch 92/100  
3500/3500 [=====] - 12s 3ms/step - loss: 17463703552.0000 - val\_loss: 15738961920.0000  
Epoch 93/100  
3500/3500 [=====] - 11s 3ms/step - loss: 17638373376.0000 - val\_loss: 15568631808.0000  
Epoch 94/100  
3500/3500 [=====] - 11s 3ms/step - loss: 17264066560.0000 - val\_loss: 15433594880.0000  
Epoch 95/100  
3500/3500 [=====] - 12s 3ms/step - loss: 17110662144.0000 - val\_loss: 15322685440.0000

```
Epoch 96/100
3500/3500 [=====] - 11s 3ms/step - loss: 17199093760.0000 - val_loss: 15301575680.0000
Epoch 97/100
3500/3500 [=====] - 12s 3ms/step - loss: 16939251712.0000 - val_loss: 15160739840.0000
Epoch 98/100
3500/3500 [=====] - 12s 3ms/step - loss: 16743148544.0000 - val_loss: 14951175168.0000
Epoch 99/100
3500/3500 [=====] - 12s 3ms/step - loss: 16887285760.0000 - val_loss: 14833034240.0000
Epoch 100/100
3500/3500 [=====] - 12s 3ms/step - loss: 16801883136.0000 - val_loss: 14721144832.0000
```

```
In [55]: pd.DataFrame({'True Values': y_test, 'Predicted Values': pred}).hvplot.scatter(x='True Values', y='Predicted Values')
```

Out[55]:

```
In [56]: pd.DataFrame(r.history)
```

Out[56]:

	loss	val_loss
0	1.635959e+12	1.658110e+12
1	1.635208e+12	1.656222e+12
2	1.630735e+12	1.647849e+12
3	1.615667e+12	1.623587e+12
4	1.578035e+12	1.568295e+12
...	...	...
95	1.719909e+10	1.530158e+10
96	1.693925e+10	1.516074e+10
97	1.674315e+10	1.495118e+10
98	1.688729e+10	1.483303e+10
99	1.680188e+10	1.472114e+10

100 rows × 2 columns

```
In [57]: pd.DataFrame(r.history).hvplot.line(y=['loss', 'val_loss'])
```

Out[57]:

```
In [59]: test_pred = model.predict(X_test)
train_pred = model.predict(X_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)

print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Artificial Neural Network", *evaluate(y_test, test_pred), 0]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])

47/47 [=====] - 0s 1ms/step
110/110 [=====] - 0s 1ms/step
Test set evaluation:

_____
MAE: 94912.90125053687
MSE: 14721153465.518364
RMSE: 121330.76059070248
R2 Square 0.8752553764041622

Train set evaluation:

_____
MAE: 98039.4612827314
MSE: 15688677933.532608
RMSE: 125254.45274932387
R2 Square 0.876923101706429
```

## Random Forest Regressor

```
In [60]: from sklearn.ensemble import RandomForestRegressor

rf_reg = RandomForestRegressor(n_estimators=1000)
rf_reg.fit(X_train, y_train)

test_pred = rf_reg.predict(X_test)
train_pred = rf_reg.predict(X_train)

print('Test set evaluation:\n_____')
print_evaluate(y_test, test_pred)

print('Train set evaluation:\n_____')
print_evaluate(y_train, train_pred)

results_df_2 = pd.DataFrame(data=[["Random Forest Regressor", *evaluate(y_test, test_pred), 0]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])
```

Test set evaluation:

---

```
MAE: 94384.86904540827
MSE: 14193674223.405785
RMSE: 119137.20755249295
R2 Square 0.8797251484003632
```

---

Train set evaluation:

---

```
MAE: 35320.769058241756
MSE: 1986369265.366525
RMSE: 44568.7027561553
R2 Square 0.9844170318823071
```

---

## Support Vector Machine

In [61]:

```
from sklearn.svm import SVR

svm_reg = SVR(kernel='rbf', C=1000000, epsilon=0.001)
svm_reg.fit(X_train, y_train)

test_pred = svm_reg.predict(X_test)
train_pred = svm_reg.predict(X_train)

print('Test set evaluation:\n' + print_evaluate(y_test, test_pred))

print('Train set evaluation:\n' + print_evaluate(y_train, train_pred))

results_df_2 = pd.DataFrame(data=[["SVM Regressor", *evaluate(y_test, test_pred), 0]],
                             columns=['Model', 'MAE', 'MSE', 'RMSE', 'R2 Square', 'Cross Validation'])
```

Test set evaluation:

---

```
MAE: 87205.73051021632
MSE: 11720932765.275507
RMSE: 108263.25676458984
R2 Square 0.9006787511983232
```

---

Train set evaluation:

---

```
MAE: 73692.56848073176
MSE: 9363827731.41128
RMSE: 96766.87310960957
R2 Square 0.9265412370487788
```

---

## Summary

We covered a lot of ground including:

- The common linear regression models (Ridge, Lasso, ElasticNet, ...).
- The representation used by the model.
- Learning algorithms used to estimate the coefficients in the model.
- Rules of thumb to consider when preparing data for use with linear regression.
- How to evaluate a linear regression model.